



BC data Workshop


BCSA

Group members:
Javier Hernandez, Alastair Jamieson-Lane, Jie
Jian, Hans Oeri, Yi Sui, Shanzhao Wang





The Data

The Data



**7000
image
files**

The Data






Challenges: Context

Challenges: Context



Challenges: Context



Challenges: Context






Challenges: Not many images

Challenges: Not many images



**7000
image
files**

Challenges: Not many images



**7000
image
files**



**14,197,122
image files**

Challenges: Not many images

VGG16

InceptionV3

ResNet

~100 million
parameters



Transfer Learning



Transfer Learning


VGG16

Transfer Learning

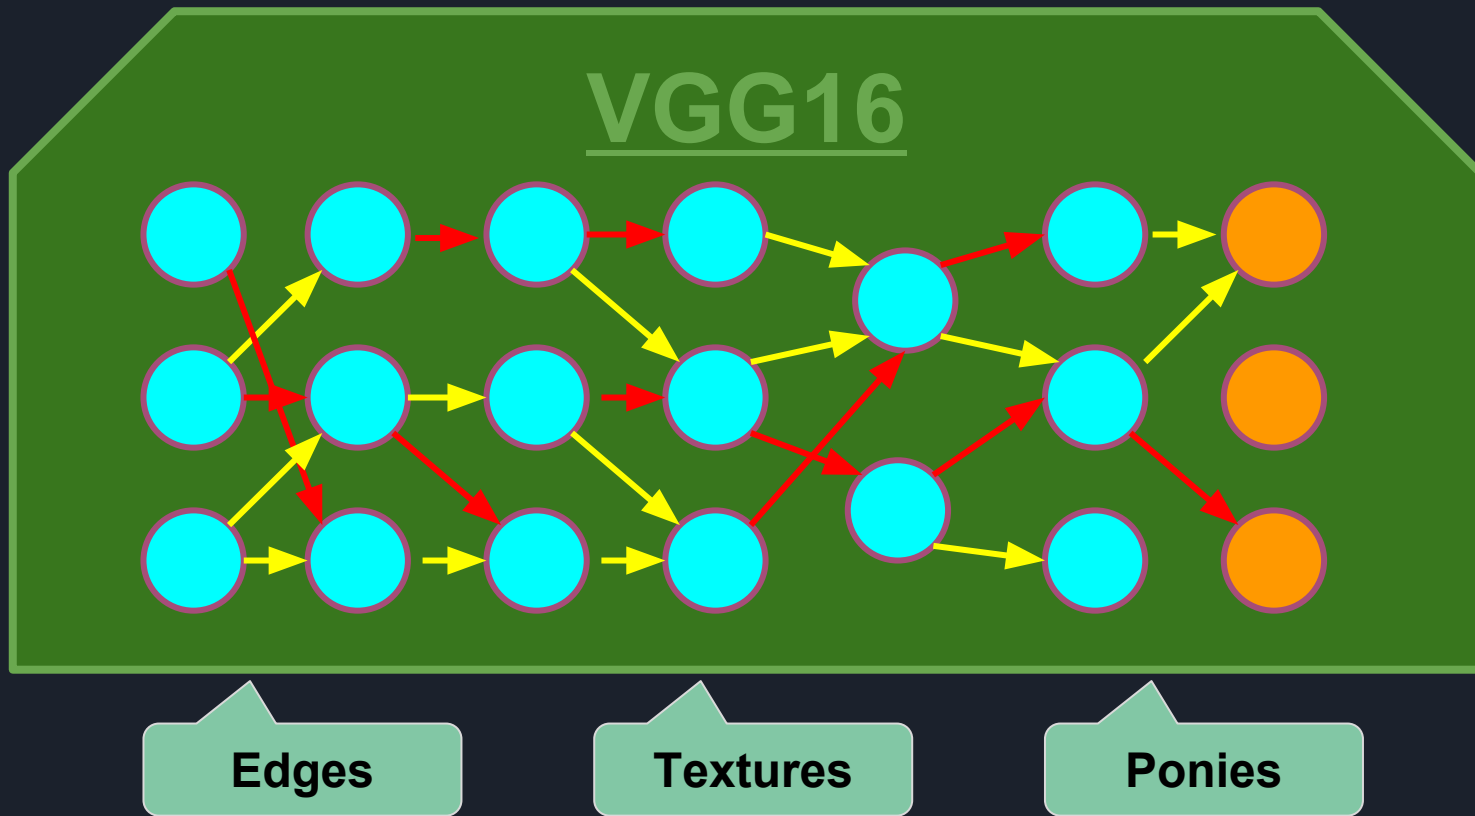


VGG16


Transfer Learning




Transfer Learning




Transfer Learning




Transfer Learning



Transfer Learning



Transfer Learning







Challenges

The data set we have to train on is relatively small for the degree of abstraction we have.

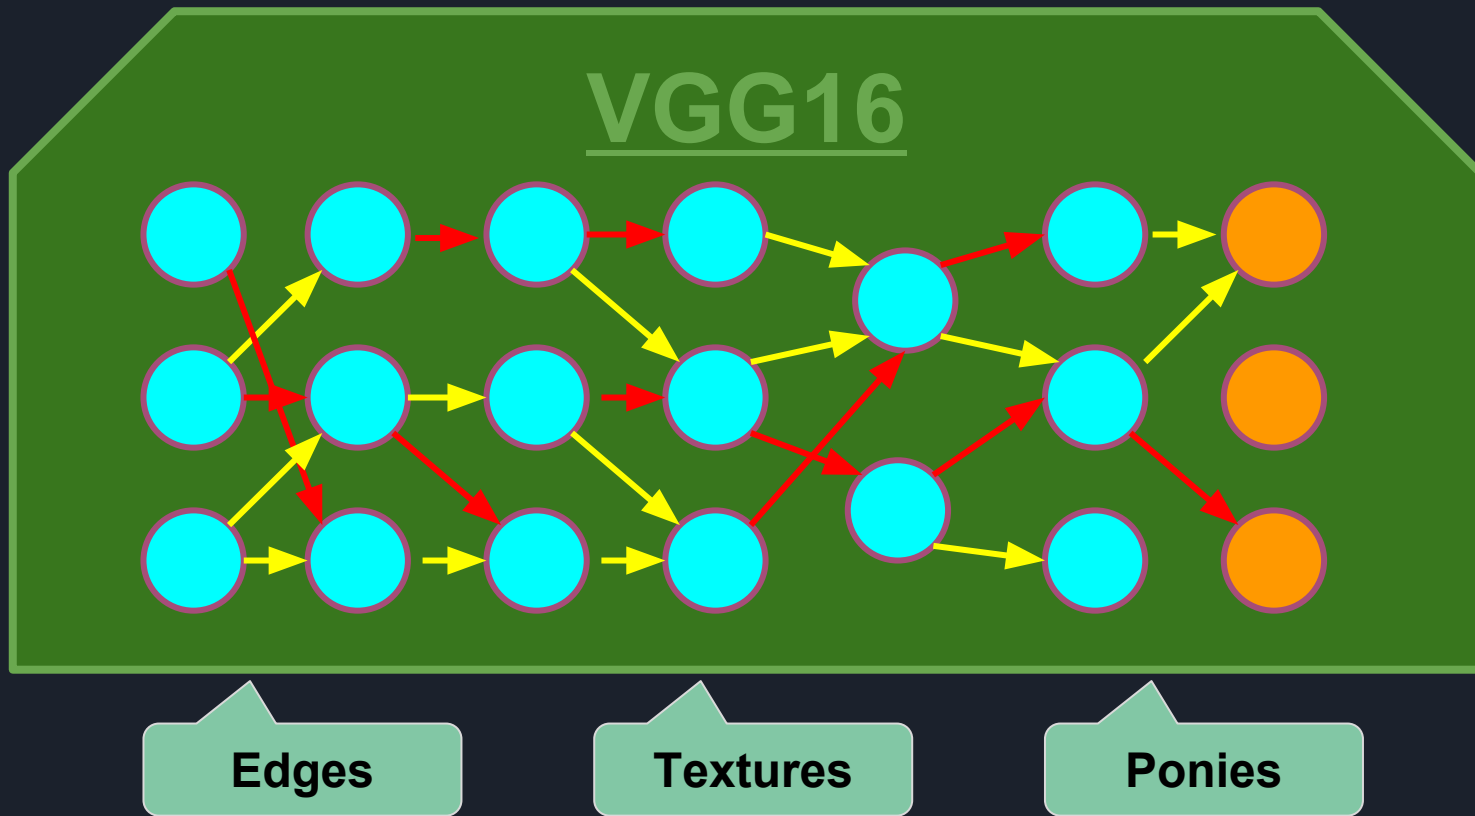
Despite small data sample, the storage needed is still very large and difficult to handle.

Training networks with many layers is time consuming.




Bottlenecking


Bottlenecking




Bottlenecking




Bottlenecking



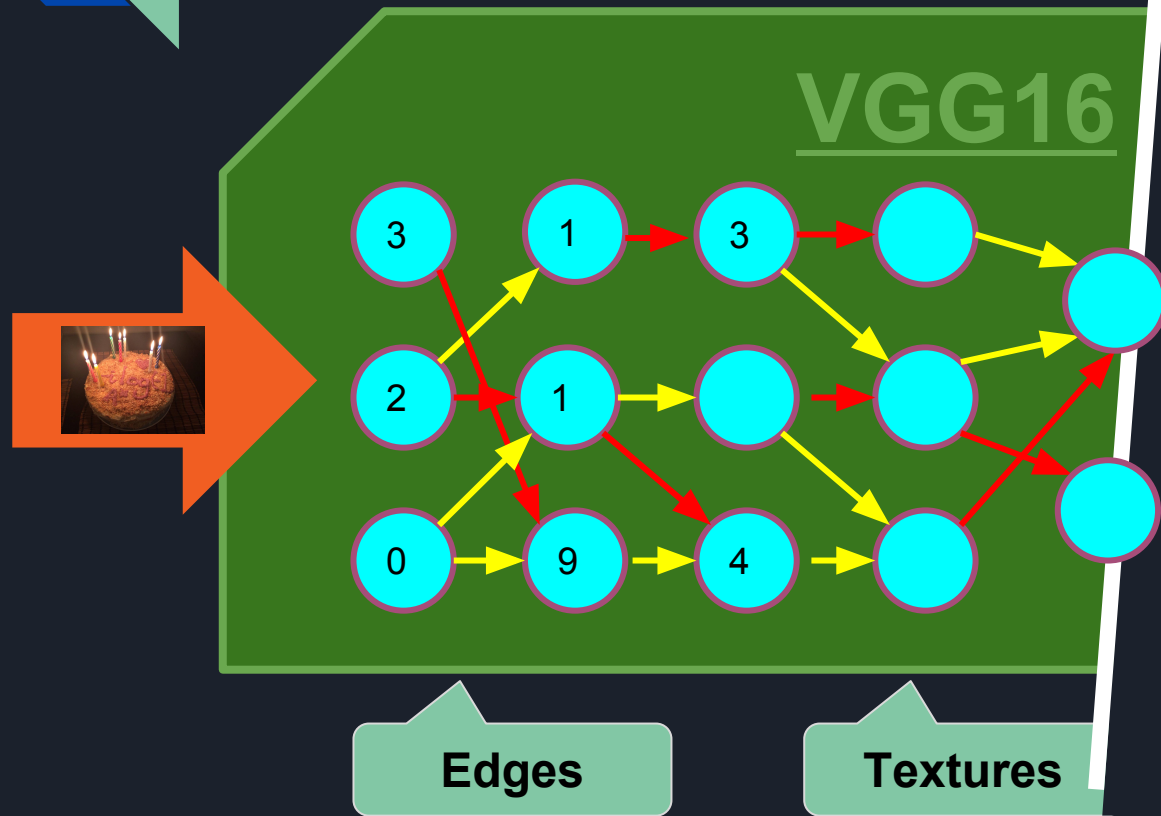
Bottlenecking



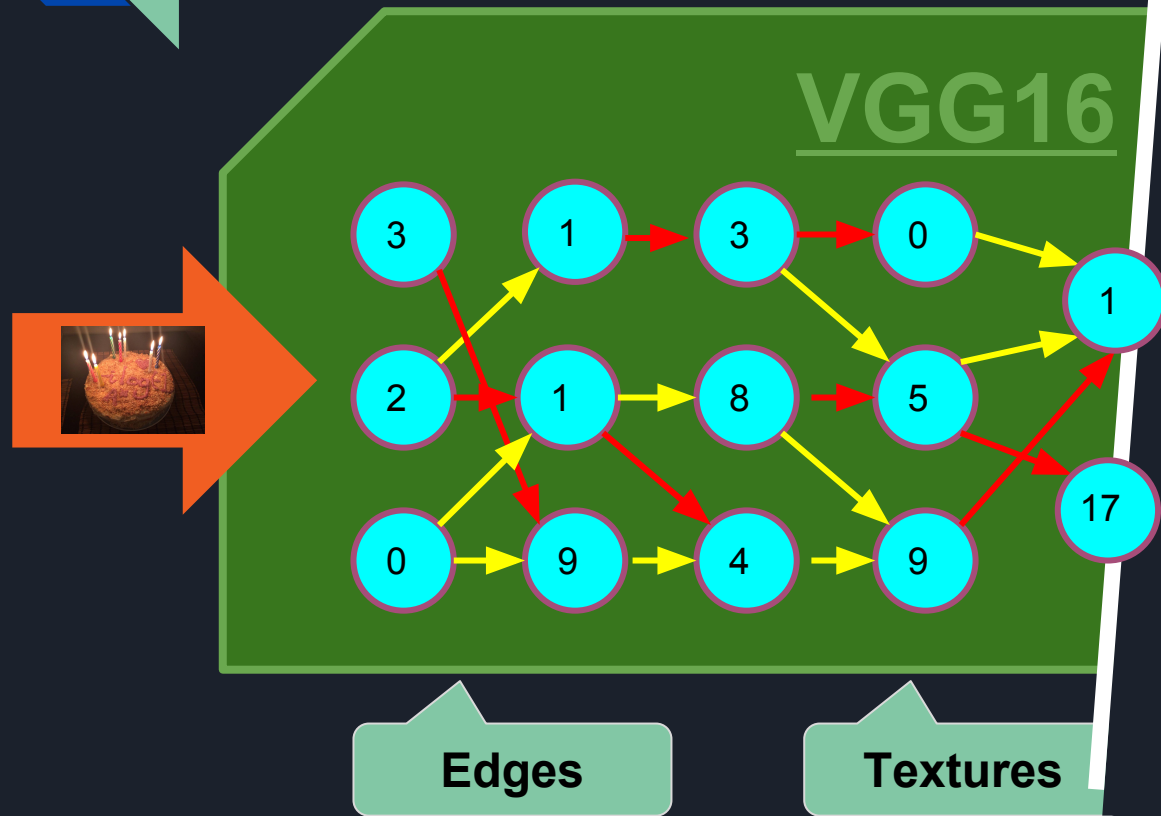
Bottlenecking




Bottlenecking




Bottlenecking




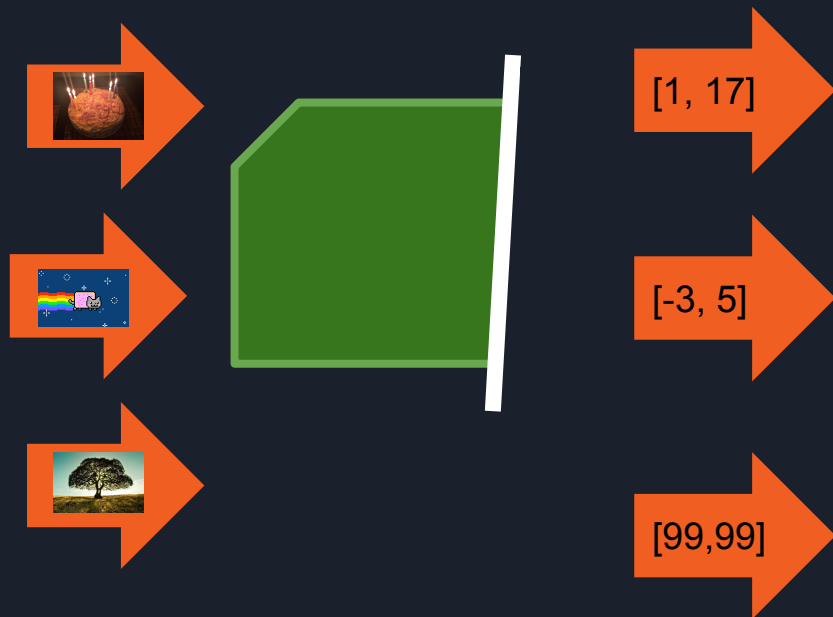
Bottlenecking



Bottlenecking



Bottlenecking




Bottlenecking

[1, 17]

[-3, 5]

[99, 99]

New Model



Bottlenecking

Weights are much smaller than image files


[1, 17]

[-3, 5]

No longer need to save image files, or VGG16 model

[99, 99]


New Model




Model is smaller than VGG16 model



Stock image set: Daisies Vs Roses



Transfer Learning (InceptionV3)



Transfer Learning (InceptionV3):

```
# create the base pre-trained model
base_model = InceptionV3(weights='imagenet', include_top=False)

# add a global spatial average pooling layer
x = base_model.output
x = GlobalAveragePooling2D()(x)
# Let's add a fully-connected layer
x = Dense(1024, activation='relu')(x)

# and a logistic layer -- For 2 classes
predictions = Dense(2, activation='softmax')(x)

# this is the model we will train
model = Model(input=base_model.input, output=predictions)

# first: train only the top layers (which were randomly initialized)
# i.e. freeze all convolutional InceptionV3 layers
for layer in base_model.layers:
    layer.trainable = False

# compile the model (should be done *after* setting layers to non-trainable)
model.compile(optimizer='rmsprop', loss='categorical_crossentropy', metrics=['accuracy'])
```



Examples (InceptionV3):



[[0.02206078 0.97793919]]



[[0.9629941 0.03700593]]

Training time:
1-2 Hours

Accuracy on
Test set: ~94%

Results (Bottlenecking)



694 images
41Mb



5 minutes

VGG16

1	73	-5	0.6
76	5	-1	2.3
3	52	18	3.5
8	1.1	5.3	9.7

1 table
12.5Mb

Results (Bottlenecking)

1	73	-5	0.6
76	5	-1	2.3
3	52	18	3.5
8	1.1	5.3	9.7

New Model

2 dense layers
1 softmax layer

Training time: 12 seconds
Accuracy on Test set: 98%



Results (Comparison)

Given the tests on the flowers data set:

- Transfer Learning: inceptionV3 Training (1-2 hours)
- Bottlenecking: VGG16 Training (12 seconds)

We chose to do bottlenecking on the BCSA data

Results (w/ BCSA Data)



7100 images
~15 GB



60 minutes

VGG16

1	78	-5	0.6
7	5	3	2
3	5	1	3
8	1	5	9

1 Data table
27MB

Results (w/ BCSA Data)



1	73	-5	0.6
76	5	-1	2.3
3	52	18	3.5
8	1.1	5.3	9.7



New Model

2 dense layers
1 softmax layer

Training time:
73 seconds

Accuracy on Test set:
~79%

Results (w/ BCSA Data)



1	73	-5	0.6
76	5	-1	2.3
3	52	18	3.5
8	1.1	5.3	9.7



New Model

2 dense layers
1 softmax layer

Training time:
73 seconds

Accuracy on Test set:
~79%

BCSA Data:
High Hazard ~80%
Low Hazard ~20%



Caveats for Bottlenecking

- Bottlenecking only works if the locked nodes are well trained. If they make wrong predictions, then the end nodes will be trained incorrectly.
- Choosing the optimal amount of top layers to train.



Further Approaches

Rather than take objects from many categories and directly trying to determine their safety, determine safety for specific type of object first. (Hazards types vary between objects)

Design the network in two steps: 1) Determine what type of object is presented. 2) then determine its safety.



Acknowledgements

We would thank BCSA for the Data and the tutors:

(Doris and Soyeon in particular)

As well as the BC Data workshop and its funders (and Aaron for organising)

And the various blogs we have stolen code from



SFU Mathematics



References

- Keras
(<https://github.com/fchollet/keras>)
- VGG16
(<https://github.com/fastai/courses/blob/master/deeplearning1/nbs/lesson1.ipynb>)
- “Using Transfer Learning and Bottlenecking to Capitalize on State of the Art DNNs”
(<https://medium.com/@galen.ballew/transferlearning-b65772083b47>)
- Flower Photos
(http://download.tensorflow.org/example_image%20%20A6/flower_photos.tgz)